

Prepared by: Bilisuma Tadesse

Product and Business Dossier

AI-Powered Mental Health Triage and Provider-Match Platform

Field	Details
Document Type	Product and Business Dossier
Project	ClearPath — AI-Powered Mental Health Triage and Navigation
Version	v1.0 — Final Draft
Date	June 4, 2026
Live Prototype	https://clearpathcare-demo.vercel.app
Companion Documents	Research Dossier; Software Requirements Specification

Document purpose. This dossier presents the complete product strategy, design rationale, and business case for ClearPath. It is intended to be read alongside the Research Dossier (evidence base and citations) and the Software Requirements Specification (architecture and technical specifications). All product decisions are traceable to research findings documented in the companion Research Dossier.

Table of Contents

- 1. Executive Summary
- 2. Problem Framing
- 3. Product Strategy: Design Thinking Process
- 4. Product Definition: What ClearPath Is and Is Not
- 5. The AI Engine: Core Differentiator
- 6. MVP1 Specification
- 7. Business Model, ROI, and Stakeholders
- 8. Roadmap Summary
- 9. Risk, Ethics, and Safety
- 10. MVP1 Rationale

1. Executive Summary

1.1 The Problem in One Sentence

In the United States, the failure of mental healthcare is not a lack of willingness to seek help — it is the system's inability to convert that intent into care quickly and correctly.

1.2 The Evidence (Summary)

The research evidence is documented fully in the companion Research Dossier. Key anchors:

- In 2024, approximately **61.5 million U.S. adults (23.4%)** had a diagnosable mental illness; approximately **48% (~29.5 million)** received no treatment in the past year [Research Dossier §3.1]
- The most-cited national wait benchmark between first outreach and a first appointment at community mental health centers is **48 days** [Research Dossier §5.1]
- A **majority of psychologists** report no openings for new patients; roughly **1 in 3** privately insured people struggle to find an in-network therapist [Research Dossier §4.3, §5.2]
- Poor employee mental health drives an estimated **\$47.6 billion/year** in absence-related losses and **\$44 billion/year** in depression-linked productivity loss [Research Dossier §6.2, §6.3]

1.3 The Product

ClearPath is a web application that uses AI as a precision triage and care-navigation engine — not a therapy chatbot. The core loop:

1. User builds a short profile (ZIP, insurance, primary concern, care format preference)
2. User completes **structured, clinically validated screening** using verbatim PHQ-9 and GAD-7 instruments
3. All scoring is **deterministic code** — the LLM never computes a clinical number
4. The system maps scores to a stepped-care recommendation and surfaces a **ranked shortlist of matched, in-network providers**
5. A plain-language, non-diagnostic summary is generated to help the user understand their results and next steps

Value proposition: Reduce the time from *"I need help"* to *"I have a shortlist of the right, reachable providers"* from weeks to minutes — safely, without making diagnostic claims.

1.4 The Business

ClearPath is sold as a **B2B SaaS mental health navigation benefit to self-insured employers** (per-employee-per-month pricing). The buyer is present (HR/Benefits at mid-market companies) and the ROI is measurable (time-to-care, productivity recovery, benefit utilization).

1.5 Current Status

MVP1 is **live and demonstrable** at <https://clearpathcare-demo.vercel.app>. The full core loop is functional and ready for live demonstration.

2. Problem Framing

2.1 Core Unmet Need

A motivated, employed adult with moderate anxiety or depression spends weeks navigating a fragmented, opaque, and unresponsive system before a first appointment — or gives up entirely — because there is no intelligent triage layer that converts willingness to seek help into matched, available care.

When someone in the U.S. decides today to seek mental health support, they encounter:

- **No structured self-assessment** — nothing helps them understand what type of care they need before searching
- **~48 days average wait** at community mental health centers between first outreach and first appointment [\[Ref. 16\]](#)
- **A majority chance** their chosen therapist is not accepting new patients [\[Ref. 20\]](#)
- **Confusing insurance landscape** — roughly 1 in 3 privately insured people report difficulty finding an in-network therapist [\[Ref. 15\]](#)
- **No stepped-care logic** — the system defaults to weekly individual therapy regardless of clinical need

The result: people never enter care, enter mismatched care, or abandon the process after weeks of frustration.

2.2 Primary User

Attribute	Description
Demographics	25–45 years old, employed, insured or employer-covered, moderate-to-high digital literacy
Clinical state	Mild-to-moderate anxiety or depression; not in acute crisis
Behavioral pattern	Has searched for help; hit walls of unavailability or cost confusion; stalled
Core need	Fast, personalized, actionable guidance — not generic information
Channel	Reachable through employer benefit programs

Secondary user (MVP2+): Employer HR/Benefits managers seeking aggregate utilization data and population-level behavioral health insight.

2.3 What the Current System Gets Wrong

Failure Mode	Why It Persists
Unstructured intake	Patients self-refer with no validated symptom data; clinicians re-collect what PHQ-9/GAD-7 capture in 5 minutes
Undirected matching	Directories match by keyword and geography — not clinical severity or specialty
Invisible wait times	Little real-time availability data; users apply to multiple providers with no feedback
Insurance opacity	Coverage often unknown until after the first visit
No stepped-care logic	Defaults to weekly therapy regardless of need; under-uses evidence-based stepped-care models
Passive crisis handling	Most tools display a hotline number; few run active, layered triage

3. Product Strategy: Design Thinking Process

ClearPath was scoped and designed using a structured **Design Thinking** methodology: Empathize, Define, Ideate, Prototype, Test.

3.1 Empathize

Methods: Desk research across federal and policy sources (CDC/NCHS, SAMHSA, HRSA, KFF, Commonwealth Fund) and peer-reviewed journals; user journey mapping from documented access barriers; competitive landscape analysis.

Key empathy findings:

1. **Shame about not knowing where to start** — the system offers no onboarding or self-assessment for new users
2. **First contact is designed for administrators, not patients** — cold-calling, paperwork, and non-responsive intake forms
3. **Employed adults are the least-reached by community centers but most reachable through employer benefits** — a commercially accessible channel with a motivated, capable user
4. **Structure reduces stigma** — users complete a validated questionnaire more readily than an open "describe your problems" field

3.2 Define

Problem statement:

A motivated, employed adult experiencing moderate anxiety or depression spends weeks navigating a fragmented, opaque system before their first appointment — or gives up — because there is no intelligent triage layer that converts their willingness to seek help into matched, available care.

How Might We statements:

- HMW reduce the time between "I need help" and "I have a first appointment" from weeks to minutes?
- HMW use validated clinical data — not free-text self-description — to improve provider match quality?
- HMW make insurance and cost information legible before a user commits to a provider?
- HMW detect escalating distress in a digital-first workflow without replacing clinical judgment?

3.3 Ideate: Three Concepts Evaluated

Concept	Core Description	AI Role
A — ClearPath (triage + match)	AI-guided profile + structured validated screening to clinical profile to ranked provider match	Navigation pacing, optional story context, crisis classifier, plain-language summary; deterministic scoring and routing
B — Bridge (between-session CBT companion)	AI delivers CBT homework between therapy sessions for existing patients	Adaptive content delivery, session summaries, progress tracking
C — PulseDesk (employer ROI dashboard)	Aggregates EAP and absence data into a mental health ROI view	Anomaly detection, NLG recommendations, predictive risk scoring

3.4 Decision Matrix

Criterion (Weight)	A: ClearPath	B: Bridge	C: PulseDesk
User impact (25%)	5 — addresses the entry barrier for ~29.5M untreated adults	4 — helps existing patients; smaller pool	3 — employer sees value; end-user impact indirect
Feasibility in 7-day window (25%)	4 — intake + match is buildable; no EHR integration needed	3 — requires provider-side integration	2 — requires data partnerships and HR integrations
Safety / risk (20%)	4 — crisis escalation well-defined; scope is pre-clinical navigation	3 — between-session AI raises liability if patient deteriorates	2 — risk scores on aggregate data create governance concerns
Differentiation (15%)	5 — no AI-first triage + match product serves the mid-market at this fidelity	3 — established AI companion players already exist	4 — adjacent enterprise products exist; ROI framing differentiated

Criterion (Weight)	A: ClearPath	B: Bridge	C: PulseDesk
Business viability (15%)	5 — employer PEPM or insurer per-referral; clear monetization path	4 — per-provider subscription; smaller market	4 — employer SaaS; long sales cycles in year 1
Weighted total	4.55 (Selected)	3.50	2.95

Decision: Concept A — ClearPath. Highest weighted score; addresses the largest underserved population; most feasible to demonstrate within the build window; clearest safety boundary and monetization path.

3.5 Prototype and Test

The working prototype (MVP1) implements the full core loop. Validation is conducted with three user personas:

- **Persona 1:** Mild anxiety (low PHQ-9/GAD-7 scores; coaching or self-guided routing)
- **Persona 2:** Moderate depression (several elevated PHQ-9 items; item 9 = 0; therapy-level routing)
- **Persona 3:** Crisis-adjacent (item 9 $\geq$ 1; crisis care level; dual-layer safety UX)

Acceptance metrics: intake completion time at or under 7 minutes, scoring accuracy on 10 test cases per instrument, 100% crisis detection rate on item 9 $\geq$ 1, provider match relevance verified manually.

4. Product Definition: What ClearPath Is and Is Not

4.1 What ClearPath Is

ClearPath is a **mental health triage and care-navigation tool**. Its primary outputs are:

1. A structured, non-diagnostic clinical profile derived from validated PHQ-9 and GAD-7 screening
2. A stepped-care recommendation (self-guided, coaching, therapy, psychiatry, urgent, or crisis) based on deterministic scoring
3. A ranked shortlist of matched, in-network providers filtered by clinical need, insurance, format preference, and location
4. A plain-language, non-diagnostic summary to support the user's decision-making

4.2 What ClearPath Is Not

Not	Why This Boundary Matters
Not a therapy chatbot	ClearPath does not provide therapy or therapeutic interventions
Not a diagnostic tool	PHQ-9 and GAD-7 are validated screening instruments; clinical diagnosis requires a licensed provider
Not a crisis management system	ClearPath applies dual-layer crisis detection and surfaces safety resources but does not replace clinical crisis services

Not	Why This Boundary Matters
Not a replacement for clinical judgment	All clinical decisions are deferred to the matched provider; the product supports navigation, not care delivery

4.3 Design Principles

1. **Demo-first:** Every feature must contribute to a working end-to-end demonstration
2. **Safety is non-negotiable:** Crisis detection and deterministic scoring are always Must-Have; never deferred
3. **Determinism over cleverness:** All clinical scoring, severity thresholds, and care-level logic are pure deterministic code
4. **Instrument validity:** PHQ-9 and GAD-7 totals come only from user-confirmed 0–3 answers — never from narrative text alone
5. **Lean architecture:** Single Next.js application; no unnecessary service layers in MVP1
6. **Cut scope, not quality:** When time is short, defer a Should-Have feature; never a Must-Have or safety requirement


5. The AI Engine: Core Differentiator

5.1 Why AI Is Core, Not Decorative

The alternative to ClearPath is a static 16-question form plus a keyword-matched directory — which already exists and systematically under-serves users. AI adds value in four ways that a form cannot:

1. **Guided journey:** Welcome copy, profile wizard support, chapter transitions, and results narrative reduce drop-off compared to a cold 16-item form
2. **Empathetic context:** Optional free-text narrative personalizes the results summary; the crisis classifier runs on that text
3. **Severity-sensitive routing:** PHQ-9 scores, GAD-7 scores, and stated concern combine into a multi-dimensional profile, enabling precision matching rather than single-keyword lookup
4. **Real-time crisis detection:** The classifier monitors for indirect crisis language in optional context, in addition to the deterministic PHQ-9 item 9 check

5.2 AI Architecture (Logical)



Core principle: The LLM never computes or reports a clinical score. Scoring, severity classification, care-level mapping, and provider matching are all deterministic code, operating on user-confirmed item answers.

5.3 AI vs. Deterministic Roles

Function	Handled by	Rationale
PHQ-9 / GAD-7 total scores	Deterministic code	Eliminates hallucination risk; auditable
Severity band mapping	Deterministic code	Clinical thresholds require exactness
Care-level recommendation	Deterministic code	Must be consistent across all runs
Provider match ranking	Deterministic code	Filter and sort logic, not generative
Story draft item suggestions	LLM (user confirms each)	Requires language understanding
Crisis classification (text)	Rules + LLM (≥0.85 confidence threshold)	Text pattern requires semantic understanding
Results summary narrative	LLM	Plain-language explanation requires generation
Welcome / chapter transitions	LLM or templated fallback	Navigation warmth

5.4 Safety Safeguards

Risk	Engineering Mitigation
LLM gives clinical advice or diagnosis	Hard system-prompt constraints; templated non-diagnostic output; adversarial testing
Crisis signal missed	Dual-layer: deterministic item-9 check + crisis-language classifier; always-visible safety card + static /crisis page
Story scoring without user confirmation	AI suggestions stored separately; scoring module uses confirmed answers only
Hallucinated score	Scoring is deterministic; LLM never handles the numeric computation
Result mistaken for diagnosis	Explicit, repeated disclaimers: "screening, not a diagnosis"
Severe symptoms conflated with crisis	Severe scores route to urgent guidance, distinct from the crisis care level and 988 resources

6. MVP1 Specification

6.1 MoSCoW Prioritization

Must-Have (All Shipped in MVP1)

ID	Feature	Rationale
M1	Deterministic PHQ-9 + GAD-7 scoring engine with unit tests	Clinical backbone; removes all hallucination risk from scores
M2	Structured PHQ-9/GAD-7 screening (verbatim items, batched UI; optional story-to-AI-draft-to-user-review path)	Validated instruments; both questionnaire and narrative entry paths
M3	Dual-layer crisis detection + safety resources (inline banner + static /crisis page)	Safety-critical; non-blocking design allows screening to complete
M4	Care-level recommendation (stepped-care mapping; urgent vs. crisis distinct)	Actionable routing without diagnostic claims
M5	Provider match engine + results page (seeded dataset)	Primary user payoff: ranked, relevant shortlist
M6	Profile onboarding wizard + landing + welcome journey	Entry point and user context collection
M7	Vercel deployment + README	Submission deliverable; live, demonstrable prototype


Should-Have

ID	Feature	Status
S1	Print/save care packet (print stylesheet on results)	Shipped
S2	LLM plain-language results narrative (non-diagnostic)	Shipped
S3	Accessibility pass (WCAG 2.1 AA) on intake + results	Target Jun 3
S4	"No matches found" fallback to community resources	Shipped
S5	Progress indicator + chapter introductions in screening	Shipped
S6	Processing pipeline UI (score, match, summary steps)	Shipped

Can Wait — MVP2 and Beyond

Real provider data pipeline, accounts and authentication, employer dashboard, provider notifications, live insurance verification, payments, FastAPI AI service, LangGraph orchestration, native mobile, multi-language support.

6.2 User Flow



Key UX principle: **no hard redirects on crisis detection**. A user who triggers crisis detection is shown the SafetySupportCard inline; the /crisis page is offered voluntarily; and the full screening flow and results remain accessible. This ensures that even users with elevated risk receive a care plan and provider shortlist alongside safety resources.

6.3 Acceptance Criteria

Criterion	Pass Condition	Status
PHQ-9 scoring accuracy	System score matches manual scoring for 10 test cases	Verified
GAD-7 scoring accuracy	System score matches manual scoring for 10 test cases	Verified
Crisis detection — item 9	PHQ-9 item 9 ≥ 1 sets isCrisis and crisis care level 100% of the time	Verified
Crisis UX	Safety resources shown inline + on results; screening not hard-blocked	Verified
Story path integrity	AI-suggested item values require user confirmation before scoring	Verified
Scoring integrity	Zero LLM-generated scores in 20 runs; scoring is deterministic	Verified
Intake completion time	Median ≤ 7 minutes across 5 runs	Verified
Match relevance	Matched providers reflect stated insurance + concern	Verified
Responsiveness	Renders correctly at 375px viewport	Verified
Accessibility	WCAG 2.1 AA for intake form elements	Target Jun 3

7. Business Model, ROI, and Stakeholders

7.1 Monetization Phases

Phase 1 (Months 1–6): Employer B2B SaaS

ClearPath as a mental health navigation benefit, offered free to employees and purchased by HR/Benefits teams.

- **Pricing model:** Per-employee-per-month (PEPM), targeting \$3–8 PEPM
- **Contract structure:** Annual contracts; minimum ~200 employees
- **Sales motion:** Direct to HR/Benefits at 500–5,000-employee companies (mid-market); avoid Fortune 500 in year 1 due to longer sales cycles

- **Why employers pay:** Poor employee mental health drives measurable costs; ClearPath reduces time-to-care, the lever most directly tied to those costs

Phase 2 (Months 6–18): Insurer/Health Plan Integration

- **Model:** PMPM or per-completed-referral fee
- **Value proposition:** Reduced avoidable ED visits, step-appropriate routing, improved parity compliance infrastructure
- **Regulatory context:** Mental health parity enforcement is evolving; payers need demonstrable structured triage [\[Ref. 13, 14\]](#)

Phase 3 (Months 18–36): Provider Network

- **Model:** Premium directory placement or per-booked-appointment referral fees
- **Governance:** Must be disclosed and structured to prevent pay-to-rank dynamics that would compromise match integrity

7.2 Unit Economics (Indicative)

Metric	Assumption	Value
Average PEPM price	Conservative	\$5/employee/month
Average contract size	500 employees	\$2,500 MRR
Year-1 target	50 employers	~\$125,000 MRR (~\$1.5M ARR)
COGS (API + hosting + support)	Estimated	~30% of revenue
Target gross margin	—	~70%
Break-even employer count	~\$30K/month operating cost	~12 employers

These are illustrative planning assumptions, not financial forecasts.

Employer ROI anchor: One 500-employee company with an average salary of \$65,000 and a 5% reduction in depression-related productivity loss recovers approximately \$163,000/year against a \$30,000/year ClearPath cost — an estimated 5:1 return, before accounting for reduced absenteeism or improved benefit utilization.

7.3 Stakeholder Map

Stakeholder	Primary Interest	Pain ClearPath Addresses	Key Concern
Patients / end users	Fast, affordable, matched care	Long waits + insurance opacity	Privacy; data misuse
Employers / HR teams	Productivity, retention, measurable benefit ROI	Impaired productivity; employees unsure how to access benefits	Liability; employee trust
Health insurers / plans	Cost containment, parity compliance, reduced ED visits	Mis-routed care; out-of-network overuse	Regulatory scrutiny

Stakeholder	Primary Interest	Pain ClearPath Addresses	Key Concern
Therapists / counselors	Full caseloads with well-matched patients	Intake burden; poor-fit referrals	Replacement concerns; data ownership
Primary care providers	Structured mental health referral pathway	No organized MH navigation for co-morbid patients	Integration complexity
Regulators (FDA, state boards)	Safety; non-diagnostic claims	Unvalidated apps making clinical claims	Perception of diagnosing vs. screening

7.4 Competitive Positioning

Category	Strength	Gap ClearPath Fills
Provider directories (Psychology Today, Zocdoc)	Large provider databases	No validated intake; no clinical-severity matching
Teletherapy marketplaces (BetterHelp, Talkspace)	Scale and brand recognition	Closed ecosystems; no external or in-person routing
Enterprise platforms (Spring Health, Lyra)	Clinically validated; enterprise-grade	Higher cost; designed for large employers; not mid-market
AI companions (Woebot, Wysa)	Clinically studied ongoing support	Support focus, not triage and navigation
Wellness apps (Headspace, Calm)	Strong B2B distribution	Wellness focus, not clinical-severity routing

ClearPath's differentiated position: AI-first triage-and-match for the **underserved mid-market** (200–5,000-employee companies), with validated clinical intake as the core engine — not wellness content, not a closed therapy ecosystem, not enterprise-only pricing.

8. Roadmap Summary

Full roadmap detail is in the companion SRS. This section summarizes the strategic phases.

8.1 MVP1: Prove the Loop (7-Day Build — Completed)

Goal: A live, demonstrable prototype proving the full core loop safely.

Delivered: Profile, structured validated screening, deterministic scoring, dual-layer crisis detection, stepped-care recommendation, ranked provider match, and results page.

Live: <https://clearpathcare-demo.vercel.app>

8.2 MVP2: Validate and Harden (Days 8–30)

Goal: Turn the prototype into a pilot-ready product with real data, real users, and the first employer conversation.

Priority	Feature
1	Real provider data pipeline (NPPEs NPI registry + public listings; ~500 providers, 5 metros)
2	Opt-in longitudinal tracking (magic-link sessions; 30-day re-screen; trend view)
3	Clinical advisory validation (2–3 licensed clinicians review intake quality and routing logic)
4	Employer dashboard v0 (aggregate-only: completion counts, care-level distribution, time-to-match)
5	First employer pilot (at least 1 signed; at least 200 employees; instrument time-to-first-appointment vs. 48-day baseline)

MVP2 success metrics:

Metric	Target
Intake completion rate	> 65%
Provider match click-through	> 35%
Crisis escalation rate (100% audited)	Low single-digit percent
Real provider coverage	≥ 500 providers, 5 metros
Employer pilots signed	≥ 1

8.3 MVP3: Close the Loop and Scale (Months 2–4)

- Provider notification and consented referral (closed-loop attribution for Phase 2 monetization)
- Live insurance eligibility verification
- Booking handoff to provider scheduling
- Embeddings-based semantic matching (complement rule-based ranking)
- FastAPI AI orchestration service (when AI work outgrows Next.js API routes)

8.4 Scale and Enterprise (Months 4+)

- LangGraph-orchestrated adaptive multi-instrument screening (PCL-5, AUDIT)
- SSO/SAML, HRIS/EAP integrations, HIPAA-compliant data architecture
- Multi-language support (Spanish first)
- Bias and fairness auditing on match outcomes
- White-label and co-brand configurations

Core principle across all phases: The batch intake and deterministic scoring infrastructure built in MVP1 is never replaced — all future work is additive.

9. Risk, Ethics, and Safety

9.1 Clinical and Safety Risks

Risk	Likelihood	Impact	Mitigation
Acute crisis not escalated	Low (dual-layer detection)	Critical	Deterministic item-9 check + parallel LLM classifier; inline SafetySupportCard + always-accessible static /crisis page
LLM gives clinical advice or diagnosis	Medium (without safeguards)	High	Hard system-prompt constraints; adversarial testing; deterministic scoring; non-diagnostic templated output
User substitutes app for professional care	Medium	High	Repeated, prominent disclaimers: "screening, not diagnosis or treatment"
Provider data inaccuracy	Medium	Medium	"Verify before booking" prompt; "last updated" timestamp on all provider cards
Model drift or prompt degradation	Low-Medium	Medium	Scheduled prompt re-evaluation; adversarial testing on every model/prompt change
Prompt injection via user input	Medium	High	Hardened system prompts; clinical decisions not LLM-driven; adversarial testing

9.2 Ethics

Scope of practice: ClearPath performs screening, not diagnosis. PHQ-9 and GAD-7 are validated screening tools used routinely in primary care worldwide. The product makes no diagnostic claims and defers clinical interpretation to the matched provider. The intended regulatory positioning is consistent with FDA's General Wellness policy; a formal determination is required prior to launch.

Algorithmic bias: Matching may underserve rural ZIPs, non-English speakers, or Medicaid enrollees due to provider data distribution. Planned mitigations: "No matches found" fallback surfacing community mental health center alternatives; language-preference filter (MVP2 roadmap); demographic-parity audits on match outcomes (MVP2 roadmap).

Data minimization: MVP1 stores no personally identifiable information. Session data does not persist beyond the browser session. No user data is sold or shared.

Informed consent: Users are informed at intake start what is collected, how it is used, that the tool is not a therapist, and that they can stop at any time.

9.3 Regulatory Posture

ClearPath is intended to operate as a **general wellness / health-navigation tool**, not Software as a Medical Device (SaMD), in alignment with FDA's General Wellness Policy [\[Ref. 54\]](#) . Key guardrails:

- No diagnostic claims are made

- PHQ-9 and GAD-7 are used as validated screening instruments, consistent with their intended use in primary care
- Clinical decisions (diagnosis, treatment planning) are deferred to the matched licensed provider
- A formal regulatory assessment must be completed before launch

If the product is later extended to include clinical decision support, FDA de novo or 510(k) clearance may be required.

HIPAA: MVP1 stores no PHI and has no covered-entity relationship. A HIPAA-compliant architecture and Business Associate Agreements are required before any employer or insurer pilot involving identified user data.

9.4 Known MVP1 Limitations

1. Provider data is seeded — real-time availability is not reflected; users are prompted to verify before booking
2. Insurance cost estimates are indicative — coverage must be confirmed with the provider
3. No independent clinical validation has been conducted — planned for MVP2 with licensed clinician advisors
4. The tool requires internet access and moderate digital literacy — digital-divide concerns are acknowledged
5. English-only in MVP1 — a meaningful equity limitation

10. MVP1 Rationale

The ClearPath MVP1 is the right starting point for seven compounding reasons:

1. Highest-leverage, most tractable failure. The intent-to-care gap can be meaningfully reduced with software today — without requiring regulatory change, new clinical partnerships, or EHR integrations. Workforce shortages and insurance reform are intractable in a 7-day window; intelligent triage navigation is not.

2. AI is structurally necessary, not decorative. A static form plus a keyword directory already exists. It under-serves. AI enables guided pacing, ambiguity handling, real-time crisis detection, and multi-dimensional matching — none achievable with a form alone.

3. The buyer is present. Employers are actively expanding mental health benefit offerings. A buyer-present, ROI-measurable market with an identified decision-maker (HR/Benefits) is the most efficient path to early revenue.

4. The safety boundary is clear and defensible. Screening is not diagnosis. Deterministic scoring removes the primary clinical risk of AI health tools. The scope is bounded, testable, and aligned with established regulatory frameworks.

5. It is shippable in the window without cutting safety corners. No EHR integration, no HIPAA-covered data, no real-time API partnerships, and no regulatory approvals are needed to prove the core loop.

6. The path to revenue is concrete. One 500-employee pilot at \$5 PEPM = \$2,500 MRR. Ten pilots = \$25,000 MRR. The B2B SaaS motion is proven in adjacent categories and directly applicable here.

7. Every future feature is additive, not a rebuild. Longitudinal tracking, booking handoff, employer dashboards, insurer integration, and adaptive multi-instrument screening all layer onto the same intake and scoring infrastructure built in MVP1. Nothing done in the 7-day build is throwaway work.